

# Solutions to G. Grolemund & H. Wickhams's R for Data Science, Chapter 5

*Krista DeStasio*

*8/2/2017*

## Contents

<b>A Brief Introduction to This File</b>	<b>2</b>
<b>Chapter 5, Data transformation</b>	<b>2</b>
Introduction . . . . .	2
<b>filter()</b> . . . . .	3
Comparisons . . . . .	4
Logical operators . . . . .	5
Missing values . . . . .	6
Exercises 5.2.4 . . . . .	7
1. Find all flights that . . . . .	7
2. Another useful dplyr filtering helper is <code>between()</code> . What does it do? Can you use it to simplify the code needed to answer the previous challenges? . . . . .	10
3. How many flights have a missing <code>dep_time</code> ? What other variables are missing? What might these rows represent? . . . . .	11
4. Why is <code>NA ^ 0</code> not missing? Why is <code>NA   TRUE</code> not missing? Why is <code>FALSE &amp; NA</code> not missing? Can you figure out the general rule? ( <code>NA * 0</code> is a tricky counterexample!) . . . . .	12
<b>arrange()</b> . . . . .	12
Exercises 5.3.1 . . . . .	13
1. How could you use <code>arrange()</code> to sort all missing values to the start? (Hint: use <code>is.na()</code> ). . . . .	13
2. Sort <code>flights</code> to find the most delayed flights. Find the flights that left earliest. . .	14
3. Sort <code>flights</code> to find the fastest flights . . . . .	14
4. Which flights travelled the longest? Which travelled the shortest? . . . . .	15
<b>select()</b> . . . . .	16
<b>rename()</b> . . . . .	17
<b>everything()</b> . . . . .	17
Exercises 5.4.1 . . . . .	18
1. Brainstorm as many ways as possible to select <code>dep_time</code> , <code>dep_delay</code> , <code>arr_time</code> , and <code>arr_delay</code> from <code>flights</code> . . . . .	18
2. What happens if you include the name of a variable multiple times in a <code>select()</code> call? .	19
3. What does the <code>one_of()</code> function do? Why might it be helpful in conjunction with this vector? . . . . .	19
4. Does the result of running the following code surprise you? How do the select helpers deal with case by default? How can you change that default? . . . . .	20
<b>mutate()</b> . . . . .	20
<b>transmute()</b> . . . . .	21
Useful creation functions . . . . .	21
Exercises 5.5.2 . . . . .	23
1. Currently <code>dep_time</code> and <code>sched_dep_time</code> are convenient to look at, but hard to compute with because they're not really continuous numbers. Convert them to a more convenient representation of number of minutes since midnight. . .	23

2. Compare `air_time` with `arr_time - dep_time`. What do you expect to see? What do you see? What do you need to do to fix it? . . . . . 24
3. Compare `dep_time`, `sched_dep_time`, and `dep_delay`. How would you expect those three numbers to be related? . . . . . 25
4. Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully read the documentation for `min_rank()`. . . . . 25
5. What does `1:3 + 1:10` return? Why? . . . . . 26
6. What trigonometric functions does R provide? . . . . . 26

## A Brief Introduction to This File

This R file walks through G. Grolemund & H. Wickhams’s online text, “R for Data Science.” Much of the code is sourced directly from the book and credit belongs to the authors. Here, some sections of code are heavily commented so that the beginning R programmer can read through and understand what each line of code does and compare it to their own as they work through the text. Throughout, the book provides the primary and most thorough explanation. **For the greatest learning benefit, I suggest you attempt each exercise on your own before looking at the code or write-ups provided here.** Of course, there is more than one way to write code and you may find a more elegant solution that you prefer.

For those new to R and RStudio, it may be of additional benefit to knit the document and examine how the code in the Rmd file is visually expressed in the resultant knitted document. For example, see how the `["R for Data Science."]` (<http://r4ds.had.co.nz/index.html>) is expressed as a hyperlink in the preceeding paragraph where it was not surrounded by tick-marks and compare that to how the same text is expressed in this paragraph when surrounded by ticks. See also the difference in appearance when knitting to different document types (HTML, PDF, Word).

**Tip:** If you are using RStudio, click the text next to the orange `#` box at the bottom of the editor window to easily navigate the code chunks.

**Tip:** Use the `?` before any command to view the documentation on that function. Do this often. For example, type `?setwd` to see a description, usage, arguments, and more for the function `setwd()`.

**Tip:** Find RStudio Cheatsheets at <https://www.rstudio.com/resources/cheatsheets/>

## Chapter 5, Data transformation

### Introduction

```
?flights
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
```

```
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
#View(flights) # This line is commented out because the View function will prevent the document from kn
```

The six key dplyr functions:

```
filter()
arrange()
select()
mutate()
summarise()
groupby()
```

## filter()

```
filter(flights, month == 1, day == 1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1     1     517           515           2     830
## 2 2013     1     1     533           529           4     850
## 3 2013     1     1     542           540           2     923
## 4 2013     1     1     544           545          -1    1004
## 5 2013     1     1     554           600          -6     812
## 6 2013     1     1     554           558          -4     740
## 7 2013     1     1     555           600          -5     913
## 8 2013     1     1     557           600          -3     709
## 9 2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
## # ... with 832 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
jan1 <- filter(flights, month == 1, day == 1)
(dec25 <- filter(flights, month == 12, day == 25))
```

```
## # A tibble: 719 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013    12    25     456           500          -4     649
## 2 2013    12    25     524           515           9     805
## 3 2013    12    25     542           540           2     832
## 4 2013    12    25     546           550          -4    1022
## 5 2013    12    25     556           600          -4     730
## 6 2013    12    25     557           600          -3     743
## 7 2013    12    25     557           600          -3     818
## 8 2013    12    25     559           600          -1     855
## 9 2013    12    25     559           600          -1     849
```

```
## 10 2013 12 25 600 600 0 850
## # ... with 709 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

## Comparisons

Comparison operators:

```
>
>=
<
<=
!=
==
```

```
# filter(flights, month = 1) # wrong
filter(flights, month == 1) # corrected
```

```
## # A tibble: 27,004 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013     1     1     517             515         2     830
## 2 2013     1     1     533             529         4     850
## 3 2013     1     1     542             540         2     923
## 4 2013     1     1     544             545        -1    1004
## 5 2013     1     1     554             600        -6     812
## 6 2013     1     1     554             558        -4     740
## 7 2013     1     1     555             600        -5     913
## 8 2013     1     1     557             600        -3     709
## 9 2013     1     1     557             600        -3     838
## 10 2013     1     1     558             600        -2     753
## # ... with 26,994 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
sqrt(2) ^ 2 == 2
```

```
## [1] FALSE
```

```
1/49 * 49 == 1
```

```
## [1] FALSE
```

```
near(sqrt(2) ^ 2, 2)
```

```
## [1] TRUE
```

```
near(1/49 * 49, 1)
```

```
## [1] TRUE
```

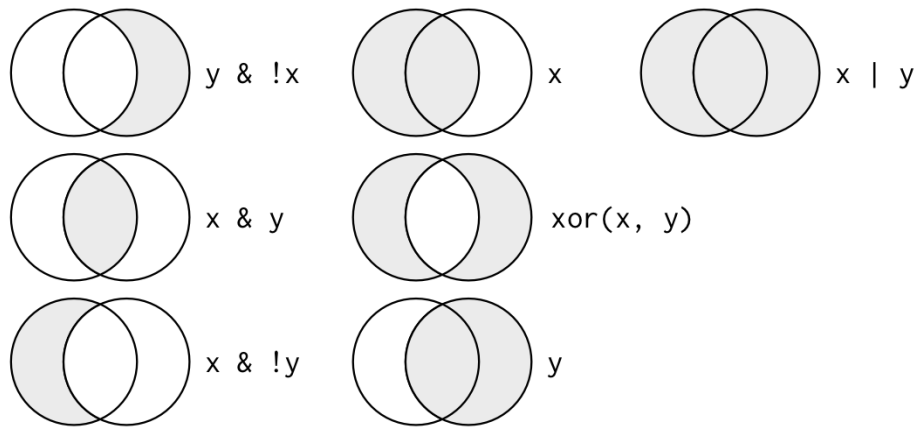


Figure 1: Boolean operations

## Logical operators

```
# Find all flights that departed in November or December
filter(flights, month == 11 | month == 12)
```

```
## # A tibble: 55,403 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013    11     1       5           2359         6       352
## 2  2013    11     1      35           2250       105       123
## 3  2013    11     1     455           500         -5       641
## 4  2013    11     1     539           545         -6       856
## 5  2013    11     1     542           545         -3       831
## 6  2013    11     1     549           600        -11       912
## 7  2013    11     1     550           600        -10       705
## 8  2013    11     1     554           600         -6       659
## 9  2013    11     1     554           600         -6       826
## 10 2013    11     1     554           600         -6       749
## # ... with 55,393 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
nov_dec <- filter(flights, month %in% c(11,12))
```

```
# Find all flights that weren't delayed (on arrival or departure) by more than two hours
filter(flights, !(arr_delay > 120 | dep_delay > 120))
```

```
## # A tibble: 316,050 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2       830
## 2  2013     1     1     533           529         4       850
## 3  2013     1     1     542           540         2       923
## 4  2013     1     1     544           545        -1      1004
## 5  2013     1     1     554           600         -6       812
## 6  2013     1     1     554           558         -4       740
```

```
## 7 2013 1 1 555 600 -5 913
## 8 2013 1 1 557 600 -3 709
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 316,040 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>

filter(flights, arr_delay <= 120, dep_delay <= 120)
```

```
## # A tibble: 316,050 x 19
##   year month day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int> <int> <int> <dbl> <int>
## 1 2013 1 1 517 515 2 830
## 2 2013 1 1 533 529 4 850
## 3 2013 1 1 542 540 2 923
## 4 2013 1 1 544 545 -1 1004
## 5 2013 1 1 554 600 -6 812
## 6 2013 1 1 554 558 -4 740
## 7 2013 1 1 555 600 -5 913
## 8 2013 1 1 557 600 -3 709
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 316,040 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

## Missing values

```
NA > 5

## [1] NA

10 == NA

## [1] NA

NA + 10

## [1] NA

NA / 2

## [1] NA

NA == NA

## [1] NA

# Let x be Mary's age. We don't know how old she is.
x <- NA

# Let y be John's age. We don't know how old he is.
y <- NA
```

```
# Are John and Mary the same age?
x == y

## [1] NA
# We don't know!

is.na(x) # Is the value of x missing

## [1] TRUE
```

## Exercises 5.2.4

### 1. Find all flights that

1. Had an arrival delay of two or more hours

```
filter(flights, arr_delay >= 120)

## # A tibble: 10,200 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     811           630        101    1047
## 2  2013     1     1     848          1835        853    1001
## 3  2013     1     1     957           733        144    1056
## 4  2013     1     1    1114           900        134    1447
## 5  2013     1     1    1505          1310        115    1638
## 6  2013     1     1    1525          1340        105    1831
## 7  2013     1     1    1549          1445         64    1912
## 8  2013     1     1    1558          1359        119    1718
## 9  2013     1     1    1732          1630         62    2028
## 10 2013     1     1    1803          1620        103    2008
## # ... with 10,190 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2. Flew to Houston (IAH or HOU)

```
filter(flights, dest %in% c("IAH", "HOU"))

## # A tibble: 9,313 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     623           627        -4     933
## 4  2013     1     1     728           732        -4    1041
## 5  2013     1     1     739           739         0    1104
## 6  2013     1     1     908           908         0    1228
## 7  2013     1     1    1028          1026         2    1350
## 8  2013     1     1    1044          1045        -1    1352
## 9  2013     1     1    1114           900        134    1447
## 10 2013     1     1    1205          1200         5    1503
## # ... with 9,303 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# or:
```

```
filter(flights, dest == "IAH" | dest == "HOU")
```

```
## # A tibble: 9,313 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     623           627        -4     933
## 4  2013     1     1     728           732        -4    1041
## 5  2013     1     1     739           739         0    1104
## 6  2013     1     1     908           908         0    1228
## 7  2013     1     1    1028          1026         2    1350
## 8  2013     1     1    1044          1045        -1    1352
## 9  2013     1     1    1114           900        134    1447
## 10 2013     1     1    1205          1200         5    1503
```

```
## # ... with 9,303 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. Were operated by United, American, or Delta

```
filter(flights, carrier %in% c("AA", "UA"))
```

```
## # A tibble: 91,394 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     554           558        -4     740
## 5  2013     1     1     558           600        -2     753
## 6  2013     1     1     558           600        -2     924
## 7  2013     1     1     558           600        -2     923
## 8  2013     1     1     559           600        -1     941
## 9  2013     1     1     559           600        -1     854
## 10 2013     1     1     606           610        -4     858
```

```
## # ... with 91,384 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
filter(flights, carrier == "AA" | carrier == "UA")
```

```
## # A tibble: 91,394 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     554           558        -4     740
## 5  2013     1     1     558           600        -2     753
## 6  2013     1     1     558           600        -2     924
```



```
## 7 2013 1 1 558 600 -2 923
## 8 2013 1 1 559 600 -1 941
## 9 2013 1 1 559 600 -1 854
## 10 2013 1 1 606 610 -4 858
## # ... with 91,384 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

4. Departed in summer (July, August, and September)

```
filter(flights, month %in% c(7, 8, 9))
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013     7     1       1         2029        212     236
## 2 2013     7     1       2         2359         3     344
## 3 2013     7     1      29         2245       104     151
## 4 2013     7     1     43         2130       193     322
## 5 2013     7     1     44         2150       174     300
## 6 2013     7     1     46         2051       235     304
## 7 2013     7     1     48         2001       287     308
## 8 2013     7     1     58         2155       183     335
## 9 2013     7     1    100         2146       194     327
## 10 2013     7     1    100         2245       135     337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
filter(flights, month == 7 | month == 8 | month == 9)
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013     7     1       1         2029        212     236
## 2 2013     7     1       2         2359         3     344
## 3 2013     7     1      29         2245       104     151
## 4 2013     7     1     43         2130       193     322
## 5 2013     7     1     44         2150       174     300
## 6 2013     7     1     46         2051       235     304
## 7 2013     7     1     48         2001       287     308
## 8 2013     7     1     58         2155       183     335
## 9 2013     7     1    100         2146       194     327
## 10 2013     7     1    100         2245       135     337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

5. Arrived more than two hours late, but didn't leave late

```
filter(flights, arr_delay > 120 & dep_delay < 1)
```

```
## # A tibble: 29 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##      <int> <int> <int>      <int>          <int>      <dbl>      <int>
## 1  2013      1    27    1419          1420        -1    1754
## 2  2013     10      7    1350          1350         0    1736
## 3  2013     10      7    1357          1359        -2    1858
## 4  2013     10     16     657           700        -3    1258
## 5  2013     11      1     658           700        -2    1329
## 6  2013      3     18    1844          1847        -3      39
## 7  2013      4     17    1635          1640        -5    2049
## 8  2013      4     18     558           600        -2    1149
## 9  2013      4     18     655           700        -5    1213
## 10 2013      5     22    1827          1830        -3    2217
## # ... with 19 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

6. Were delayed by at least an hour, but made up over 30 minutes in flight

```
filter(flights, dep_delay >= 60 & arr_delay < -30)
```

```
## # A tibble: 0 x 19
## # ... with 19 variables: year <int>, month <int>, day <int>,
## #   dep_time <int>, sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

7. Departed between midnight and 6am (inclusive)

```
filter(flights, dep_time > 0 & dep_time < 600)
```

```
## # A tibble: 8,730 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # ... with 8,720 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2. Another useful dplyr filtering helper is `between()`. What does it do? Can you use it to simplify the code needed to answer the previous challenges?

```
?between
```

```
# 1.5 simplified
filter(flights, between(flights$month, 7, 9))
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     7     1       1           2029         212     236
## 2  2013     7     1       2           2359          3     344
## 3  2013     7     1      29           2245        104     151
## 4  2013     7     1     43           2130        193     322
## 5  2013     7     1     44           2150        174     300
## 6  2013     7     1     46           2051        235     304
## 7  2013     7     1     48           2001        287     308
## 8  2013     7     1     58           2155        183     335
## 9  2013     7     1    100           2146        194     327
## 10 2013     7     1    100           2245        135     337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

*# 1.7 simplified*

```
filter(flights, between(flights$dep_time, 0, 600))
```

```
## # A tibble: 9,344 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1    517           515          2     830
## 2  2013     1     1    533           529          4     850
## 3  2013     1     1    542           540          2     923
## 4  2013     1     1    544           545         -1    1004
## 5  2013     1     1    554           600         -6     812
## 6  2013     1     1    554           558         -4     740
## 7  2013     1     1    555           600         -5     913
## 8  2013     1     1    557           600         -3     709
## 9  2013     1     1    557           600         -3     838
## 10 2013     1     1    558           600         -2     753
## # ... with 9,334 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. How many flights have a missing `dep_time`? What other variables are missing? What might these rows represent?

```
filter(flights, is.na(dep_time))
```

```
## # A tibble: 8,255 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     NA           1630         NA     NA
## 2  2013     1     1     NA           1935         NA     NA
## 3  2013     1     1     NA           1500         NA     NA
## 4  2013     1     1     NA           600          NA     NA
## 5  2013     1     2     NA           1540         NA     NA
## 6  2013     1     2     NA           1620         NA     NA
## 7  2013     1     2     NA           1355         NA     NA
## 8  2013     1     2     NA           1420         NA     NA
```

```
## 9 2013 1 2 NA 1321 NA NA
## 10 2013 1 2 NA 1545 NA NA
## # ... with 8,245 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# or
table(is.na(flights$dep_time))
```

```
##
## FALSE TRUE
## 328521 8255
```

4. Why is  $NA^0$  not missing? Why is  $NA \mid TRUE$  not missing? Why is  $FALSE \& NA$  not missing? Can you figure out the general rule? ( $NA * 0$  is a tricky counterexample!)

- $NA^0$  is not missing because any value to the 0th power is 1
- $NA \mid TRUE$  is not missing because the  $\mid$  operand will return TRUE as long as one condition is true. TRUE is TRUE.
- $FALSE \& NA$  is not missing because the NA is ignored
- In operations, any value interacting with an NA becomes missing. Missing values are ignored in conditional expressions.

`arrange()`

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013     1     1     517             515         2       830
## 2 2013     1     1     533             529         4       850
## 3 2013     1     1     542             540         2       923
## 4 2013     1     1     544             545        -1      1004
## 5 2013     1     1     554             600        -6       812
## 6 2013     1     1     554             558        -4       740
## 7 2013     1     1     555             600        -5       913
## 8 2013     1     1     557             600        -3       709
## 9 2013     1     1     557             600        -3       838
## 10 2013     1     1     558             600        -2       753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
```

```
## 1 2013 1 9 641 900 1301 1242
## 2 2013 6 15 1432 1935 1137 1607
## 3 2013 1 10 1121 1635 1126 1239
## 4 2013 9 20 1139 1845 1014 1457
## 5 2013 7 22 845 1600 1005 1044
## 6 2013 4 10 1100 1900 960 1342
## 7 2013 3 17 2321 810 911 135
## 8 2013 7 22 2257 759 898 121
## 9 2013 12 5 756 1700 896 1058
## 10 2013 5 3 1133 2055 878 1250
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>

df <- tibble(x = c(5, 2, NA))
arrange(df, x)
```

```
## # A tibble: 3 x 1
##       x
##   <dbl>
## 1     2
## 2     5
## 3    NA
```

```
arrange(df, desc(x))
```

```
## # A tibble: 3 x 1
##       x
##   <dbl>
## 1     5
## 2     2
## 3    NA
```

### Exercises 5.3.1

1. How could you use `arrange()` to sort all missing values to the start? (Hint: use `is.na()`).

```
# Example using arrival times
arrange(flights, !is.na(arr_time))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1     1    2016           1930           46     NA
## 2 2013     1     1      NA           1630           NA     NA
## 3 2013     1     1      NA           1935           NA     NA
## 4 2013     1     1      NA           1500           NA     NA
## 5 2013     1     1      NA            600           NA     NA
## 6 2013     1     2    2041           2045          -4     NA
## 7 2013     1     2    2145           2129           16     NA
## 8 2013     1     2      NA           1540           NA     NA
## 9 2013     1     2      NA           1620           NA     NA
## 10 2013     1     2      NA           1355           NA     NA
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

## 2. Sort flights to find the most delayed flights. Find the flights that left earliest.

```
# Most delayed at departure
```

```
arrange(flights, desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9     641             900         1301    1242
## 2  2013     6    15    1432            1935         1137    1607
## 3  2013     1    10    1121            1635         1126    1239
## 4  2013     9    20    1139            1845         1014    1457
## 5  2013     7    22     845            1600         1005    1044
## 6  2013     4    10    1100            1900          960    1342
## 7  2013     3    17    2321             810          911     135
## 8  2013     6    27     959            1900          899    1236
## 9  2013     7    22    2257             759          898     121
##10  2013    12     5     756            1700          896    1058
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# Left earliest
```

```
arrange(flights, dep_delay)
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013    12     7    2040            2123          -43     40
## 2  2013     2     3    2022            2055          -33    2240
## 3  2013    11    10    1408            1440          -32    1549
## 4  2013     1    11    1900            1930          -30    2233
## 5  2013     1    29    1703            1730          -27    1947
## 6  2013     8     9     729             755          -26    1002
## 7  2013    10    23    1907            1932          -25    2143
## 8  2013     3    30    2030            2055          -25    2213
## 9  2013     3     2    1431            1455          -24    1601
##10  2013     5     5     934             958          -24    1225
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

## 3. Sort flights to find the fastest flights

```
arrange(flights, desc(distance / air_time))
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1 2013 5 25 1709 1700 9 1923
## 2 2013 7 2 1558 1513 45 1745
## 3 2013 5 13 2040 2025 15 2225
## 4 2013 3 23 1914 1910 4 2045
## 5 2013 1 12 1559 1600 -1 1849
## 6 2013 11 17 650 655 -5 1059
## 7 2013 2 21 2355 2358 -3 412
## 8 2013 11 17 759 800 -1 1212
## 9 2013 11 16 2003 1925 38 17
## 10 2013 11 16 2349 2359 -10 402
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

#### 4. Which flights travelled the longest? Which travelled the shortest?

```
# Flights with the longest travel time
arrange(flights, desc(air_time))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     3    17   1337           1335          2    1937
## 2 2013     2     6    853           900         -7    1542
## 3 2013     3    15   1001           1000          1    1551
## 4 2013     3    17   1006           1000          6    1607
## 5 2013     3    16   1001           1000          1    1544
## 6 2013     2     5    900           900          0    1555
## 7 2013    11    12    936           930          6    1630
## 8 2013     3    14    958           1000         -2    1542
## 9 2013    11    20   1006           1000          6    1639
## 10 2013     3    15   1342           1335          7    1924
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# Flights with the shortest travel time
arrange(flights, air_time)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1 2013     1    16   1355           1315         40    1442
## 2 2013     4    13    537           527         10     622
## 3 2013    12     6    922           851         31    1021
## 4 2013     2     3   2153           2129         24    2247
## 5 2013     2     5   1303           1315        -12    1342
## 6 2013     2    12   2123           2130         -7    2211
## 7 2013     3     2   1450           1500        -10    1547
## 8 2013     3     8   2026           1935         51    2131
## 9 2013     3    18   1456           1329         87    1533
## 10 2013     3    19   2226           2145         41    2305
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

`select()`

```
# Select columns by name
```

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
##10  2013     1     1
## # ... with 336,766 more rows
```

```
# Select all columns between year and day (inclusive)
```

```
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
##10  2013     1     1
## # ... with 336,766 more rows
```

```
# Select all columns except those from year to day
```

```
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
##   <int>         <int>         <dbl>   <int>         <int>         <dbl>
## 1     517           515           2     830           819           11
## 2     533           529           4     850           830           20
## 3     542           540           2     923           850           33
## 4     544           545          -1    1004          1022          -18
## 5     554           600          -6     812           837           -25
## 6     554           558          -4     740           728           12
```



```
## 7      555      600      -5      913      854      19
## 8      557      600      -3      709      723     -14
## 9      557      600      -3      838      846     -8
## 10     558      600      -2      753      745      8
## # ... with 336,766 more rows, and 10 more variables: carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Helper functions to use with `select()`

- `starts_with("abc")`: matches names that begin with “abc”.
- `ends_with("xyz")`: matches names that end with “xyz”.
- `contains("ijk")`: matches names that contain “ijk”.
- `matches("(.)\\1")`: selects variables that match a regular expression. This one matches any variables that contain repeated characters. You’ll learn more about regular expressions in strings.
- `num_range("x", 1:3)`: matches `x1`, `x2` and `x3`.

`rename()`

```
rename(flights, tail_num = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tail_num <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

`everything()`

```
# Move handful of variables to the start of the data frame
select(flights, time_hour, air_time, everything())
```

```
## # A tibble: 336,776 x 19
##           time_hour air_time year month   day dep_time sched_dep_time
##           <dtm>     <dbl> <int> <int> <int>   <int>         <int>
## 1 2013-01-01 05:00:00    227  2013     1     1     517             515
## 2 2013-01-01 05:00:00    227  2013     1     1     533             529
## 3 2013-01-01 05:00:00    160  2013     1     1     542             540
## 4 2013-01-01 05:00:00    183  2013     1     1     544             545
## 5 2013-01-01 06:00:00    116  2013     1     1     554             600
## 6 2013-01-01 05:00:00    150  2013     1     1     554             558
```

```
## 7 2013-01-01 06:00:00      158 2013      1      1      555      600
## 8 2013-01-01 06:00:00       53 2013      1      1      557      600
## 9 2013-01-01 06:00:00      140 2013      1      1      557      600
## 10 2013-01-01 06:00:00     138 2013      1      1      558      600
## # ... with 336,766 more rows, and 12 more variables: dep_delay <dbl>,
## #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance <dbl>,
## #   hour <dbl>, minute <dbl>
```

## Exercises 5.4.1

1. Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

```
# select
select(flights, dep_time, dep_delay, arr_time, arr_delay)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>     <dbl>
## 1     517         2     830         11
## 2     533         4     850         20
## 3     542         2     923         33
## 4     544        -1    1004        -18
## 5     554        -6     812        -25
## 6     554        -4     740         12
## 7     555        -5     913         19
## 8     557        -3     709        -14
## 9     557        -3     838         -8
## 10    558        -2     753          8
## # ... with 336,766 more rows
```

```
# starts_with
select(flights, starts_with("dep_"), starts_with("arr_"))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>     <dbl>
## 1     517         2     830         11
## 2     533         4     850         20
## 3     542         2     923         33
## 4     544        -1    1004        -18
## 5     554        -6     812        -25
## 6     554        -4     740         12
## 7     555        -5     913         19
## 8     557        -3     709        -14
## 9     557        -3     838         -8
## 10    558        -2     753          8
## # ... with 336,766 more rows
```

```
# ends_with
select(flights, ends_with("time"), ends_with("delay"), -(starts_with("sched")), -(starts_with("air")))
```

```
## # A tibble: 336,776 x 4
##   dep_time arr_time dep_delay arr_delay
##   <int>     <int>     <dbl>     <dbl>
```

```
## 1      517      830      2      11
## 2      533      850      4      20
## 3      542      923      2      33
## 4      544     1004     -1     -18
## 5      554      812     -6     -25
## 6      554      740     -4      12
## 7      555      913     -5      19
## 8      557      709     -3     -14
## 9      557      838     -3      -8
## 10     558      753     -2       8
## # ... with 336,766 more rows
```

```
# contains
select(flights, contains("dep"), contains("arr_"), -(contains("sched")))
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_delay arr_time arr_delay
##   <int>     <dbl>   <int>     <dbl>
## 1      517         2      830         11
## 2      533         4      850         20
## 3      542         2      923         33
## 4      544        -1     1004        -18
## 5      554        -6      812        -25
## 6      554        -4      740         12
## 7      555        -5      913         19
## 8      557        -3      709        -14
## 9      557        -3      838         -8
## 10     558        -2      753          8
## # ... with 336,766 more rows
```

## 2. What happens if you include the name of a variable multiple times in a select() call?

The variable is included only once in the data frame

```
select(flights, arr_delay, dep_delay, dep_delay)
```

```
## # A tibble: 336,776 x 2
##   arr_delay dep_delay
##   <dbl>     <dbl>
## 1      11         2
## 2      20         4
## 3      33         2
## 4     -18        -1
## 5     -25        -6
## 6      12        -4
## 7      19        -5
## 8     -14        -3
## 9      -8        -3
## 10      8        -2
## # ... with 336,766 more rows
```

## 3. What does the one\_of() function do? Why might it be helpful in conjunction with this vector?

```
vars <- c("year", "month", "day", "dep_delay", "arr_delay")
```

`one_of()` allows one to select variables that match those in a character vector.

```
vars <- c("year", "month", "day", "dep_delay", "arr_delay")
select(flights, one_of(vars))
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2         11
## 2  2013     1     1         4         20
## 3  2013     1     1         2         33
## 4  2013     1     1        -1        -18
## 5  2013     1     1        -6        -25
## 6  2013     1     1        -4         12
## 7  2013     1     1        -5         19
## 8  2013     1     1        -3        -14
## 9  2013     1     1        -3         -8
## 10 2013     1     1        -2          8
## # ... with 336,766 more rows
```

4. Does the result of running the following code surprise you? How do the select helpers deal with case by default? How can you change that default?

```
select(flights, contains("TIME"))
```

The `select()` helpers are not case sensitive by default and will select all variables that contain the character string `time` regardless of case (e.g. `"TIME"`, `"Time"`, `time`, `"TiMe"`, etc.). To change the default, set `ignore.case = FALSE` within the helper function.

```
# Default
select(flights, contains("TIME"))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time
##   <int>         <int>     <int>         <int>     <dbl>
## 1     517           515       830           819       227
## 2     533           529       850           830       227
## 3     542           540       923           850       160
## 4     544           545      1004          1022       183
## 5     554           600       812           837       116
## 6     554           558       740           728       150
## 7     555           600       913           854       158
## 8     557           600       709           723         53
## 9     557           600       838           846       140
## 10    558           600       753           745       138
## # ... with 336,766 more rows, and 1 more variables: time_hour <dtm>
```

```
# Set ignore.case = FALSE
select(flights, contains("TIME", ignore.case = FALSE))
```

```
## # A tibble: 336,776 x 0
```

`mutate()`

```
# Create a narrower dataset
flights_sml <- select(flights, year:day, ends_with("delay"), distance, air_time)

# Create new columns
mutate(flights_sml,
       gain = arr_delay - dep_delay,
       speed = distance / air_time * 60,
       hours = air_time / 60,
       gain_per_hour = gain / hours)

## # A tibble: 336,776 x 11
##   year month   day dep_delay arr_delay distance air_time  gain  speed
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl> <dbl> <dbl>
## 1  2013     1     1         2       11    1400    227     9 370.0441
## 2  2013     1     1         4       20    1416    227    16 374.2731
## 3  2013     1     1         2       33    1089    160    31 408.3750
## 4  2013     1     1        -1      -18    1576    183   -17 516.7213
## 5  2013     1     1        -6      -25     762    116   -19 394.1379
## 6  2013     1     1        -4       12     719    150    16 287.6000
## 7  2013     1     1        -5       19    1065    158    24 404.4304
## 8  2013     1     1        -3      -14     229     53   -11 259.2453
## 9  2013     1     1        -3       -8     944    140    -5 404.5714
## 10 2013     1     1        -2        8     733    138    10 318.6957
## # ... with 336,766 more rows, and 2 more variables: hours <dbl>,
## #   gain_per_hour <dbl>
```

```
transmute()
```

Only keep the new variables.

```
transmute(flights,
          gain = arr_delay - dep_delay,
          hours = air_time / 60,
          gain_per_hour = gain / hours)
```

```
## # A tibble: 336,776 x 3
##   gain      hours gain_per_hour
##   <dbl>   <dbl>   <dbl>
## 1     9 3.7833333     2.378855
## 2    16 3.7833333     4.229075
## 3    31 2.6666667    11.625000
## 4   -17 3.0500000    -5.573770
## 5   -19 1.9333333    -9.827586
## 6    16 2.5000000     6.400000
## 7    24 2.6333333     9.113924
## 8   -11 0.8833333    -12.452830
## 9    -5 2.3333333     -2.142857
## 10   10 2.3000000     4.347826
## # ... with 336,766 more rows
```

## Useful creation functions

1. Modular arithmetic, including:

- Division

```
100 / 3
```

```
## [1] 33.33333
```

- Integer division

```
100 %/% 3
```

```
## [1] 33
```

- Remainder division

```
100 %% 3
```

```
## [1] 1
```

*# Breaking up is(n't) hard to do*

```
transmute(flights,
  dep_time,
  hour = dep_time %/%100,
  minute = dep_time %% 100)
```

```
## # A tibble: 336,776 x 3
```

```
##   dep_time  hour minute
```

```
##   <int> <dbl> <dbl>
```

```
## 1     517     5     17
```

```
## 2     533     5     33
```

```
## 3     542     5     42
```

```
## 4     544     5     44
```

```
## 5     554     5     54
```

```
## 6     554     5     54
```

```
## 7     555     5     55
```

```
## 8     557     5     57
```

```
## 9     557     5     57
```

```
## 10    558     5     58
```

```
## # ... with 336,766 more rows
```

2. logs

3. Offsets (lead & lag)

```
(x <- 1:10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
lag(x)
```

```
## [1] NA 1 2 3 4 5 6 7 8 9
```

```
lead(x)
```

```
## [1] 2 3 4 5 6 7 8 9 10 NA
```

4. Cumulative and rolling aggregates

- cumsum()

- cumprod()

- cummin()

- cummax()

```

• cummean()

x
## [1] 1 2 3 4 5 6 7 8 9 10
cumsum(x)
## [1] 1 3 6 10 15 21 28 36 45 55
cummean(x)
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5

5. Logical comparisons
6. Ranking
(y <- c(1, 2, 2, NA, 3, 4))

## [1] 1 2 2 NA 3 4
min_rank(y) # Default: smallest values get small ranks

## [1] 1 2 2 NA 4 5
min_rank(desc(y)) # Give the largest values the smallest ranks

## [1] 5 3 3 NA 2 1
row_number(y)

## [1] 1 2 3 NA 4 5
dense_rank(y)

## [1] 1 2 2 NA 3 4
percent_rank(y)

## [1] 0.00 0.25 0.25 NA 0.75 1.00
cume_dist(y)

## [1] 0.2 0.6 0.6 NA 0.8 1.0

```

## Exercises 5.5.2

1. Currently `dep_time` and `sched_dep_time` are convenient to look at, but hard to compute with because they're not really continuous numbers. Convert them to a more convenient representation of number of minutes since midnight.

```

transmute(flights,
  dep_time,
  dep_minutes = dep_time %/% 100 * 60 + dep_time %% 100,
  sched_dep_time,
  sched_dep_minutes = sched_dep_time %/% 100 * 60 + sched_dep_time %% 100)

## # A tibble: 336,776 x 4
##   dep_time dep_minutes sched_dep_time sched_dep_minutes
##   <int>      <dbl>      <int>      <dbl>
## 1     517        317         515         315
## 2     533        333         529         329

```

```
## 3      542      342      540      340
## 4      544      344      545      345
## 5      554      354      600      360
## 6      554      354      558      358
## 7      555      355      600      360
## 8      557      357      600      360
## 9      557      357      600      360
## 10     558      358      600      360
## # ... with 336,766 more rows
```

Alternatively, write a function to convert time to minutes as seen here: <https://jrnold.github.io/e4qf/data-transformation.html#mutate>

```
time2mins <- function(x) {
  x %/% 100 * 60 + x %% 100
}

transmute(flights,
  dep_minutes = time2mins(dep_time),
  sched_dep_minutes = time2mins(sched_dep_time))
```

```
## # A tibble: 336,776 x 2
##   dep_minutes sched_dep_minutes
##   <dbl>         <dbl>
## 1      317         315
## 2      333         329
## 3      342         340
## 4      344         345
## 5      354         360
## 6      354         358
## 7      355         360
## 8      357         360
## 9      357         360
## 10     358         360
## # ... with 336,766 more rows
```

**2. Compare `air_time` with `arr_time - dep_time`. What do you expect to see? What do you see? What do you need to do to fix it?**

What we want is for `air_time` to equal `arr_time - dep_time`. However, since the times are not continuous values, we will get a meaningless value. We must first convert `arr_time` and `dep_time` to a format such as minutes that represents the elapsed time on a continuous scale. However, since the plane may depart from and arrive in different time zones, and since arrival and departure times are reported in local time, some of the calculated values will be different than the `air_time` values.

```
# Without converting to minutes, note that the columns do not match
transmute(flights,
  air_time,
  calculated_airtime = arr_time - dep_time)
```

```
## # A tibble: 336,776 x 2
##   air_time calculated_airtime
##   <dbl>         <int>
## 1     227         313
## 2     227         317
## 3     160         381
```



```
## 4      183      460
## 5      116      258
## 6      150      186
## 7      158      358
## 8       53      152
## 9      140      281
## 10     138      195
## # ... with 336,766 more rows

# Now that we have converted the time format, the calculation will be correct. But wait, some still don't
(flights2 <- transmute(flights,
  arr_time = time2mins(arr_time),
  dep_time = time2mins(dep_time),
  air_time,
  airtime_new = arr_time - dep_time))
```

```
## # A tibble: 336,776 x 4
##   arr_time dep_time air_time airtime_new
##   <dbl>   <dbl>   <dbl>     <dbl>
## 1     510     317     227       193
## 2     530     333     227       197
## 3     563     342     160       221
## 4     604     344     183       260
## 5     492     354     116       138
## 6     460     354     150       106
## 7     553     355     158       198
## 8     429     357      53        72
## 9     518     357     140       161
## 10    473     358     138       115
## # ... with 336,766 more rows
```

To fix this, let's simply view those flights that did not change time zones.

```
flights2 %>% filter(air_time == airtime_new)
```

```
## # A tibble: 196 x 4
##   arr_time dep_time air_time airtime_new
##   <dbl>   <dbl>   <dbl>     <dbl>
## 1     821     710     111       111
## 2    1338    1187     151       151
## 3     554     411     143       143
## 4     666     427     239       239
## 5     731     583     148       148
## 6    1106     979     127       127
## 7     754     616     138       138
## 8     974     859     115       115
## 9    1427    1315     112       112
## 10   1245    1133     112       112
## # ... with 186 more rows
```

3. Compare `dep_time`, `sched_dep_time`, and `dep_delay`. How would you expect those three numbers to be related?

4. Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully read the documentation for `min_rank()`.

5. What does `1:3 + 1:10` return? Why?
6. What trigonometric functions does `R` provide?